



Dataframes & Spark SQL



Spark SQL

Spark module
for
Structured Data Processing

Spark SQL

Integrated

- Provides **DataFrames**
- Mix SQL queries & Spark programs

Spark SQL

Uniform Data Access

- Source:
 - HDFS,
 - Hive
 - Relational Databases
- Avro, Parquet, ORC, JSON
- You can even join data across these sources.
- Hive Compatibility
- Standard Connectivity

DataFrames

RDD

1 sandeep
2 ted
3 thomas
4 priya
5 kush

Unstructured

Need code for processing

DataFrames

RDD

1	sandeep
2	ted
3	thomas
4	priya
5	kush

Unstructured

Need code for processing

Data Frame

ID	Name
1	sandeep
2	ted
3	thomas
4	priya
5	kush

Structured

Can use SQL or R like syntax:
df.sql("select Id where name = 'priya'")


head(where(df, df\$ID > 2))

Data Frames

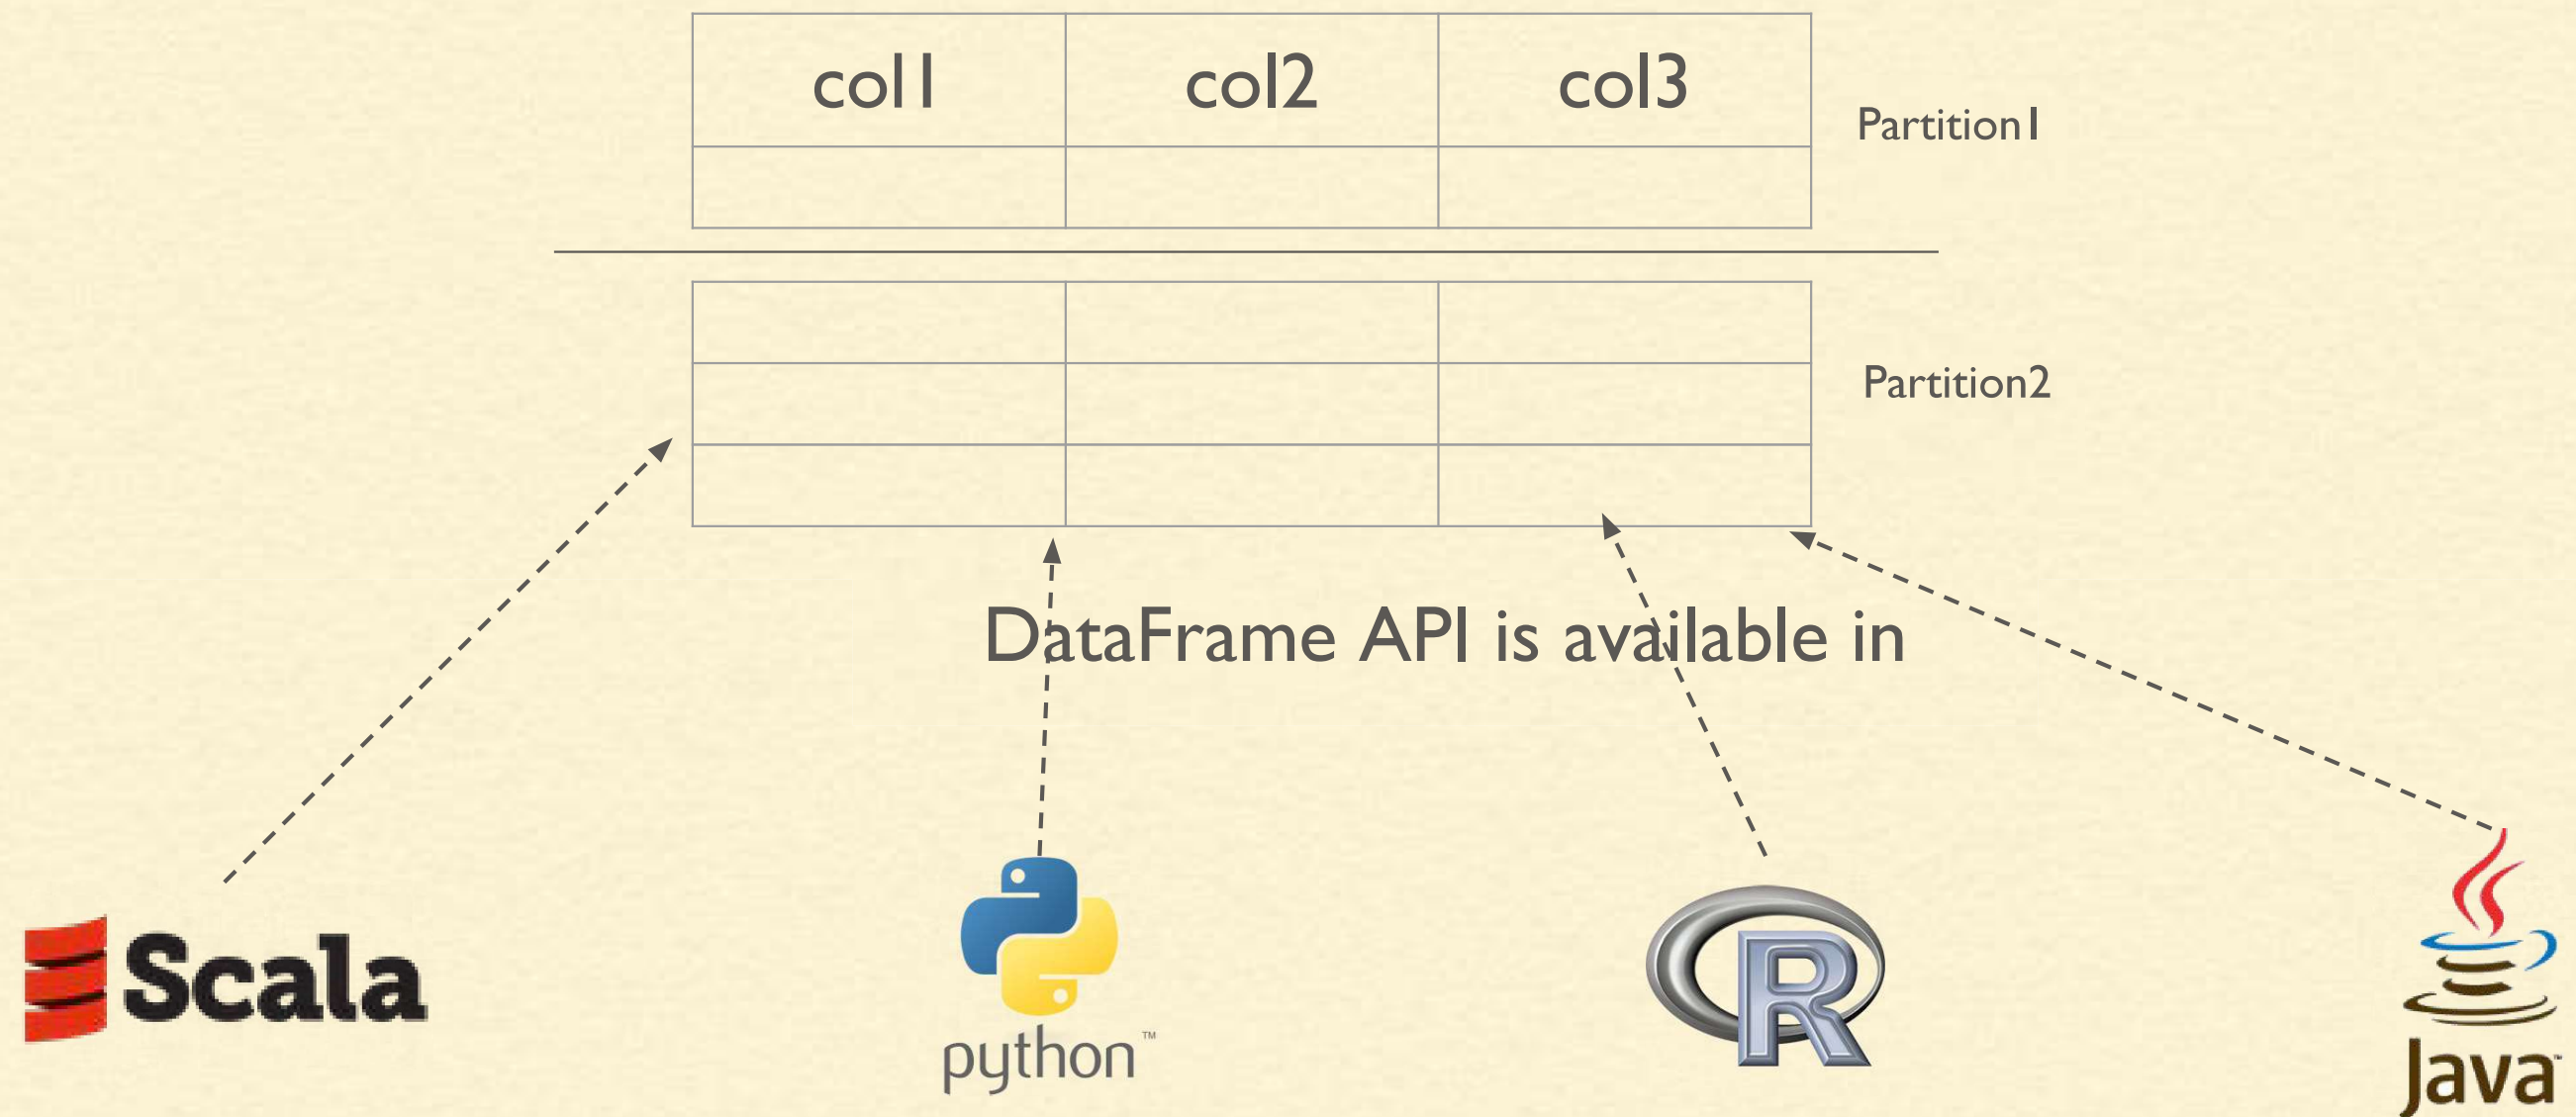
col1	col2	col3	Partition1
			Partition2

- Collection with named columns
- Distributed
- <> Same as database table
- <> A data frame in R/Python

Data Frames



Data Frames



Getting Started

- Available in Spark 2.0x onwards.
- Using usual interfaces
 - Spark-shell
 - Spark Application
 - Pyspark
 - Java
 - etc.

Getting Started

```
$ export HADOOP_CONF_DIR=/etc/hadoop/conf/  
$ export YARN_CONF_DIR=/etc/hadoop/conf/
```

Getting Started

```
$ export HADOOP_CONF_DIR=/etc/hadoop/conf/
```

```
$ export YARN_CONF_DIR=/etc/hadoop/conf/
```

```
$ ls /usr/
```

```
bin  games  include  jdk64  lib64  local  share  spark1.6  
spark2.0.2  tmp  etc  hdp  java  lib  libexec  sbin  spark1.2.1  
spark2.0.1  src
```


Getting Started

```
$ export HADOOP_CONF_DIR=/etc/hadoop/conf/
$ export YARN_CONF_DIR=/etc/hadoop/conf/
$ ls /usr/
bin  games  include  jdk64  lib64    local  share      spark1.6
spark2.0.2  tmp  etc  hdp    java    lib    libexec  sbin    spark1.2.1
spark2.0.1  src
$ /usr/spark2.0.2/bin/spark-shell
```

Getting Started

Spark context Web UI available at <http://172.31.60.179:4040>

Spark context available as 'sc' (master = local[*], app id = local-1498489557917).

```
Spark session available as 'spark'.
```

Welcome to

```

      /_/_/
     /_/_/
    /_/_/
   /_/_/
  /_/_/
 /_/_/
/_/_/

version 2.0.2

```

```
Using Scala version 2.11.8 (OpenJDK 64-Bit Server VM, Java 1.8.0_91)
```

Type in expressions to have them evaluated.

Type :help for more information.

```
scala>
```

Getting Started

Spark context Web UI available at <http://172.31.60.179:4040>

```
Spark context available as 'sc' (master = local[*], app id = local-1498489557917).
```

```
Spark session available as 'spark'.
```

Welcome to

```

      /_/_/
     /_/_/
    /_/_/
   /_/_/
  /_/_/
 /_/_/
/_/_/

version 2.0.2

```

```
Using Scala version 2.11.8 (OpenJDK 64-Bit Server VM, Java 1.8.0_91)
```

Type in expressions to have them evaluated.

Type :help for more information.

```
scala>
```

Starting Point: SparkSession

```
import org.apache.spark.sql.SparkSession

val spark = SparkSession
    .builder()
    .appName("Spark SQL basic example")
    .config("spark.some.config.option", "some-value")
    .getOrCreate()
```


Starting Point: SparkSession

```
import org.apache.spark.sql.SparkSession

val spark = SparkSession
    .builder()
    .appName("Spark SQL basic example")
    .config("spark.some.config.option", "some-value")
    .getOrCreate()

//For implicit conversions, e.g. RDDs to DataFrames
import spark.implicits._
```

Creating DataFrames from JSON

In web console or ssh:

```
$ hadoop fs -cat /data/spark/people.json
```

```
{"name": "Michael"}
```

```
{"name": "Andy", "age": 30}
```

```
{"name": "Justin", "age": 19}
```

Creating DataFrames from JSON

```
var df = spark.read.json("/data/spark/people.json")

# Displays the content of the DataFrame to stdout
df.show()
```

```
scala> df.show()
```

```
+-----+-----+
|  age|   name|
+-----+-----+
| null|Michael|
|   30|   Andy|
|   19| Justin|
+-----+-----+
```

Creating DataFrames from JSON

```
var df = spark.read.json("/data/spark/people.json")

# Displays the content of the DataFrame to stdout
df.show()
```

```
scala> df.show()
```

```
+-----+-----+
|  age |   name |
+-----+-----+
| null | Michael |
|   30 |    Andy |
|   19 |   Justin |
+-----+-----+
```



Original JSON:

```
{"name": "Michael"}
{"name": "Andy", "age": 30}
{"name": "Justin", "age": 19}
```


DataFrame Operations

```
# Print the schema in a tree format  
df.printSchema()
```

```
root  
|-- age: long (nullable = true)  
|-- name: string (nullable = true)
```

```
{"name": "Michael"}  
{"name": "Andy", "age": 30}  
{"name": "Justin", "age": 19}
```

DataFrame Operations

```
# Select only the "name" column
```

```
df.select("name").show()
```

```
+-----+  
|   name |  
+-----+  
| Michael |  
|   Andy  |  
|  Justin |  
+-----+
```

```
{"name": "Michael"}  
{"name": "Andy", "age": 30}  
{"name": "Justin", "age": 19}
```

DataFrame Operations

```
# Increment the age by 1
df.select($"name", $"age" + 1).show()
```

```
+-----+-----+
|  name | (age + 1) |
+-----+-----+
|Michael|      null |
|   Andy|       31 |
| Justin|       20 |
+-----+-----+
```

```
{"name": "Michael"}
{"name": "Andy", "age": 30}
{"name": "Justin", "age": 19}
```

DataFrame Operations

```
# Select people older than 21
df.filter($"age"> 21).show()
```

```
+---+-----+
|age|name|
+---+-----+
| 30|Andy|
+---+-----+
```

```
{"name":"Michael"}
{"name":"Andy", "age":30}
{"name":"Justin", "age":19}
```


DataFrame Operations

```
# Count people by age
df.groupBy("age").count().show()
```

```
+-----+-----+
|  age|count|
+-----+-----+
|   19|    1|
| null|    1|
|   30|    1|
+-----+-----+
```

#SQL Equivalent

```
Select age, count(*) from df group by age
```

```
{"name": "Michael"}
{"name": "Andy", "age": 30}
{"name": "Justin", "age": 19}
```

Running SQL Queries Programmatically

Running SQL Queries Programmatically

```
// Register the DataFrame as a SQL temporary view  
df.createOrReplaceTempView("people")
```

Running SQL Queries Programmatically

```
// Register the DataFrame as a SQL temporary view  
df.createOrReplaceTempView("people")
```

```
val sqlDF = spark.sql("SELECT * FROM people")
```


Running SQL Queries Programmatically

```
// Register the DataFrame as a SQL temporary view  
df.createOrReplaceTempView("people")
```

```
val sqlDF = spark.sql("SELECT * FROM people")  
sqlDF.show()
```

```
+-----+-----+  
|  age|   name|  
+-----+-----+  
|null|Michael|  
|  30|   Andy|  
|  19| Justin|  
+-----+-----+
```